

Entrepôts de Données et Big-Data

Intervenants

- Anne-Muriel Chifolleau – UM, LIRMM - Resp.
(chifolleau@lirmm.fr)
- Federico Ulliana – UM, LIRMM, INRIA - Resp.
(ulliana@lirmm.fr)
- Christophe Menichetti – IBM



Database Evolution History

1960s

First Computerized Database Models



1970s

The Dawn of the Database

- The relational model and its language SQL emerge
- The disruptive model causes the demise of other models




1970 E.F. Codd Writes a Paper on the Relational Database Model

System R by IBM (SQL)
Ingres (QUEL)

1980s

An Industry Develops


- SQL becomes the de-facto standard
- Commercial offerings from IBM, Oracle grow market
- Other data models enter the scene, without much traction



1990s

Technology Shifts

- Data explodes with the Internet age
- Single server SQL databases run into resource problems
- Business Intelligence and Analytics move out of transactional database



2000s

New Players Emerge

- Data variety, velocity and volume increase
- New analytics SQL databases are introduced
- NoSQL databases fill the gap for processing unstructured data
- Hadoop gains traction for analyzing petabytes of data

Today

Databases Adapt and Evolve

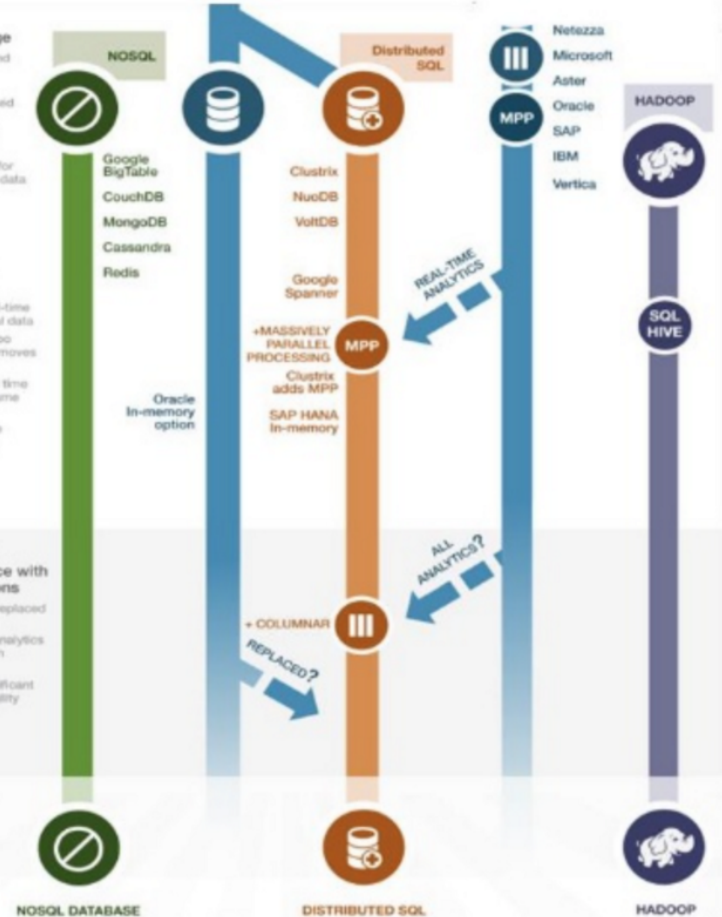
- Businesses require real-time analytics on operational data
- Scale-up SQL proves too costly, but scale-out removes resource constraint
- Scale-out provides real time analytics with high volume transactions
- Google and Clustrix are pioneers in this space

The Future

Businesses Advance with Database Innovations

- Single node SQL gets replaced by scale-out SQL
- Data warehouse type analytics will become available in real-time database
- Businesses gain a significant edge and increased agility

Winning Database Platforms



Programme

1. Rappels et préliminaires (stockage, optimisation)
2. Entrepôts de données
3. Technologies du Big-Data

Objectif : présenter les techniques de modélisation et les technologies conçues pour l'interrogation des données massives

MCC 100%CCI

(2 sessions)

- Première session
 - Rappels (rendu facultatif)
 - Optimisation (TP à rendre)
 - Entrepôts de données (mini-projet + oral)
 - Hadoop Map/Reduce (mini-projet + oral)

TP et mini-projets en binômes - oraux individuels

- Deuxième session :
 - mini-projet individuel + oral sur tout le programme
- **Présence obligatoire** aux séminaires IBM

Divers

- Consultez la page Moodle dédiée à l'UE
- TP : éléments de correction dispensés en cours
- Access comptes Oracle *possible à distance* (SSH)
 - instructions sur Moodle
- Signaler les dysfonctionnements des serveurs Oracle
 - ENT --> Assistance --> Centre de Services --> Déclarer un Incident
- Nous écrire pour tout type de problème ou question
 - chifolleau@lirmm.fr , ulliana@lirmm.fr

Rappels

Readings

These slides should be considered simply as “pointers” to the references below.

- [BD-G] Bases de données,
Georges Gardarin, 5ème edition 2005
http://georges.gardarin.free.fr/Livre_BD_Contenu/XX-TotalBD.pdf
- [ORA] Oracle® Database SQL Language
Reference 11g Release 1 (11.1) - 2013
http://docs.oracle.com/cd/B28359_01/server.111/b28286.pdf
- [UML] Prolegomenes_uml.pdf
- [UML2] UML2 : de l'apprentissage à la pratique

Relational Databases & UML

NB : Assumed to be well known from L2/L3, we just recall basic topics.

1. UML

2. Relational Model


3. SQL

Levels of Modelling

- Conceptual Model (UML, EA, Merise)
 - defines what the system contains
- Logical Model (Relational Model, Object Model, Graph)
 - Defines data structures and rules of the system
- Physical Model (SQL, OQL, XML)
 - defines how the system has to be implemented



Peter Chang
EA – 1976



Ted Codd
RM – 1970



Don Chamberlin
SQL – 1974

BASIC RELATIONAL THEORY

The Relational Model

[BD-G] chapter VI section 2.1

Everything is a relation

- **Person(Bob, 42, Paris)**
(can model entities)
- **LiveTogether(Alice, Bob, Lyon, 2010)**
(can model associations)

Relational Schema

[BD-G] chapter VI section 2.2

- A set of relations built on a set of attributes, with well defined domains.

Person(Name, Age, City)

Name : String

Age: Integer

City : {Lyon, Paris}

The Model VS. The Content

- The idea of representing data using relations is clearly independent from the data to store.
- But, as this data is originated from real world interactions (eg., trading, social), all forms of weak and strong correlations are found in it.

Functional Dependency

[BD-G] chapter VI section 3

*A set of attributes **A** determining a set of attributes **B***

(name, surname) -----> birthday
determine

city -----> (population, state)
determine

Key(s)

[BD-G] chapter VI section 3.1

minimal set of attributes determining a whole tuple

LiveTogether(Person1, Person2, City, Date)

key Person1 , Person2 , Date -----> City
 determines

(ex) **LiveTogether**(Alice , Bob , Lyon , 2010)

Strongly recommended in systems (efficiency, coherence)

Data dependencies were undesirable

- Except for **keys** and referential **integrity constraints**
 - beside these cases, they just bring redundancy
- Database **normalization** eliminated dependencies

Normal Forms

Normal-Forms are guidelines for modeling.

Their definition is motivated by design mistakes.

So, let's find the right place for the attributes !

Normal Forms : 2NF

[BD-G] chapter VI section 6.2

- **2NF** : non-key attributes fully-dependent from the key

FournisseurPiece(Name, Article, Address, Price) **NO**

(Article --/--> Address)

↓ (decomposition)

FournisseurPiece(Name, Article, Price)

Fournisseur(Name, Address)

Normal Forms : 3NF

[BD-G] chapter VI section 6.3

- 3NF : no dependencies between non-key attributes

Person(ID , Name, City, *CityPopulation*)

NO

ID -----> City -----> CityPopulation

↓ (decomposition)

Person(ID , Name, City) **Place**(City, *CityPopulation*)

Normal Forms : 3NF

[BD-G] chapter VI section 6.3

- This normal form is respected by **most** “*transactional-database*” you will find in **any** real world company
 - it allows to fix common data redundancy problems
 - also, every schema can be normalized in 3NF

Normal Forms : Remarks

- 2NF & 3NF respected in practically any information system using a relational database
- Stronger normal forms (BCNF, 4NF, 5NF) are less employed (avoid rarer mistakes; not always achievable)
- Exceptions to normalizations are tolerated to save joins (at the price of redundancy)
 - **we will see this for datawarehouses**

SQL : SURVIVAL KIT

SQL

- Structured Query Language
- Declarative (logical) Language : tell what you want from relations, not what to do with them.
 - This is the main difference with C and Java, *not only* the fact that we deal with data.
- In SQL terminology, a relation is called “table”.

SQL

- DDL (Data Definition Language)
 - CREATE/ALTER structures (table, view, index)
- DML (Data Manipulation Language)
 - UPDATE/INSERT/DELETE content
- DQL (Data Query Language)
 - SELECT data

Create Table

[BD-G] chapter VII section 2.1 and [ORA] section 16-6

```
CREATE TABLE Employee (  
  
    id    NUMBER,  
    name  VARCHAR2(50),  
    birthday DATE  
  
)
```

Oracle Built-in Datatypes

[ORA] section 2-6

Why do we need datatypes ?

- To associate fixed set of properties to attributes.
- This improves the database:
 - coherence : type-checking operations
 - cannot sum two strings
 - efficiency : a datatype has its own best storage
 - BLOB vs integers

Oracle Built-in Datatypes

- Character
- Numeric
- Date/Time
- Large Object

Complete list : see [ORA] table 2-1

Value Constraints

[ORA] section 8-4 and [BD-G] chapter VII section 2.1

- Why do we need constraints ?
- To restrict the values in a database and ensure the data integrity
 - ex : No employee without an ID

Not Null

[ORA] section 8-8

- Prohibits a database value from being null

```
CREATE TABLE Employee (  
  
  id          NUMBER NOT NULL,  
  name        VARCHAR2(50),  
  birthday    DATE  
  
)
```

Unique

[ORA] section 8-9

- Prohibits multiple rows from having the same value (but allows them to be null)

```
CREATE TABLE Employee (  
  
    id            NUMBER UNIQUE,  
    name          VARCHAR2(50),  
    birthday      DATE  
  
)
```

Primary Key

[ORA] section 8-9

- Combines a NOT NULL constraint and a UNIQUE constraint in a single declaration

```
CREATE TABLE Employee (  
  
    id            NUMBER PRIMARY KEY,  
    name          VARCHAR2(50),  
    birthday      DATE  
  
)
```


Primary Key : Multiple Attributes

[ORA] section 8-20

- Combines a NOT NULL constraint and a unique constraint in a single declaration

```
CREATE TABLE Employee (  
  id          NUMBER,  
  name        VARCHAR2(50),  
  birthday    DATE,  
  
  PRIMARY KEY (name,birthday)  
)
```

Foreign key

*one (or more) **attributes** which correspond to the **key** of another relation*

Employee(ID, Name, **Department_id)**

Departement(Dept_ID, Name)



Foreign Key

[ORA] section 8-10 and 8-21

Requires values in one table to match values in another table.

```
CREATE TABLE Employee (  
    id          NUMBER,  
    department  NUMBER  
  
    FOREIGN KEY department  
        REFERENCES Dept (dept_id)  
)
```

Check

[ORA] section 8-10 and 8-22

Requires a value to satisfy with a specified condition

```
CREATE TABLE Employee (  
    id          NUMBER,  
    department  NUMBER,  
    office      VARCHAR2 ( 10 )  
  
    CHECK (  
        office IN  
            ( ' DALLAS ' , ' BOSTON ' , ' PARIS ' , ' TOKYO ' )  
    )  
)
```

Oracle does not verify mutually exclusive conditions (eg. AGE>1 AND AGE<0)

ALTER TABLE

[BD-G] chapter VII section 6.2.4 and [ORA] section 12-2

- Add a new column
 - ALTER TABLE Employee ADD (office VARCHAR2(20));
- Modify an existing column
 - ALTER TABLE Employee MODIFY (office NUMBER);
- Define a default value for the new column
 - ALTER TABLE Employee MODIFY office DEFAULT 'Corridor';
- Drop a column
 - ALTER TABLE Employee DROP (office);

DELETE

TRUNCATE

DROP

[ORA] section 17-25 and 19-62 and 18-5

removes

rows

table

rollback

DELETE



TRUNCATE



DROP



- TRUNCATE = DROP + CREATE TABLE

INSERT

[BD-G] chapter VII section 4.1 [ORA] section 18-66 and 18-54

```
INSERT
INTO Employee
VALUES
    ( ' Bob ' ,
      TO_DATE(
        ' 03-OCT-1972 ' , ' DD-MON-YYYY ' )
    )
```

TO DATE converts a character/numeric to a date
[ORA] 2-49/50

SELECT FROM

[BD-G] chapter VII section 3.1 [ORA] section 18-66 and 18-54 and 2-49

SELECT

TO_CHAR(birthday, 'MM-DD-YYYY')

FROM

Employee

TO CHAR converts a numeric to a character

[ORA] 2-287/288 and 5-292

SELECT FROM WHERE

[BD-G] chapter VII section 3.2 [ORA] section 2-50

SELECT

name

FROM

Employee

WHERE

birthday >

TO_DATE('01-10-1970', 'DD-MM-YYYY')

JOINS

[BD-G] chapter VII section 3.3

SELECT

Employee.name, Department.name

FROM

Employee, Department

WHERE

Employee.dept = Department.id

JOINS

Employee

name	dept
Alice	dep1
Bob	dep2
Eddy	dep1

Department

id	name
dep1	Sales
dep2	Production

Emp_Join_Dep

Employee.name	Department.name
Alice	Sales
Bob	Production
Eddy	Sales

Group By

[BD-G] chapter VII section 3.7

```
SELECT
    dept, count(*) as N
FROM
    Employee
GROUP By
    dept
```

Employee

name	dept
Alice	dep1
Bob	dep2
Eddy	dep1

Agg_Emp

dept	N
dep1	2
dep2	1

Summing Up

- Relations
 - Functional dependencies, Normal-forms
- SQL
 - CREATE, INSERT, DELETE, SELECT, GROUP-BY

MODELING WITH UML

UML (Unified Modeling Language)

- Universal graphical modeling language designed to model objects, associations, time events, system states
 - Main goal is to ease prototyping
- Good news : UML is a rich model and we can also use it to model data !


UML : Plan

- UML Class Diagram
 - Basic constructs that can be used to model Relational Databases in UML
 - Real Object-oriented features

Class

[UML] section 3 [BD-G] chapter XVII section 2


- Set of elements sharing common properties
 - ex. peoples, animals, cars
- UML : draw a labelled box



Class : Attributes

[UML] section 3 [BD-G] chapter XVII section 2

- Attributes denote the properties of class objects
 - They are usually typed
- Write the attributes below the class name




Operations can specify

- Visibility (+ public) (- private) (# protected)
- Return-type (*optional*; can be undefined)
- Multiplicity of parameter/return-type (*optional*)

Employee
name : string dept : int
+setDepartment(int dept [1]) : bool [1]

Instances

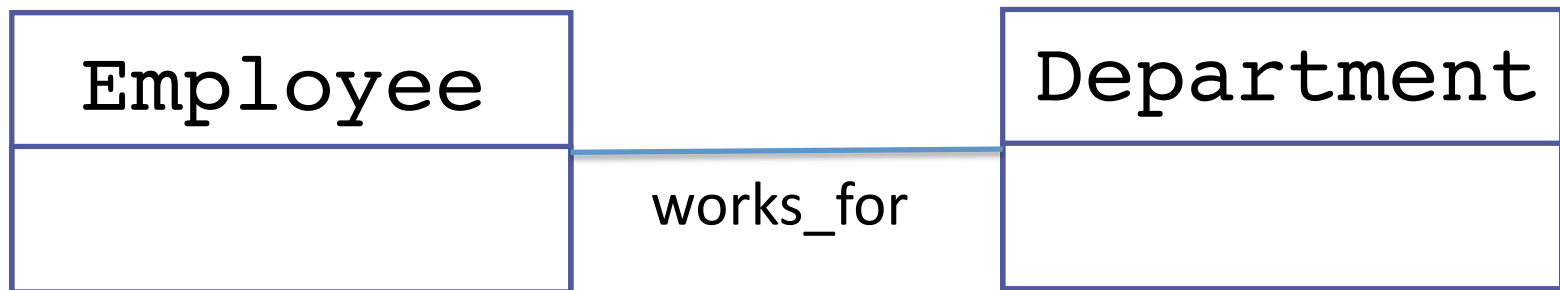
- The elements of a class
 - ex. the employee Bob
- Related to their class by a directed edge



Binary Associations


[UML] section 4 [BD-G] chapter XVII section 2

- General relationships between elements of two classes
 - ex. an employee works for a department
 - a concrete instance of association is called link
- Binary association : *undirected* edge between classes



Associations : Links

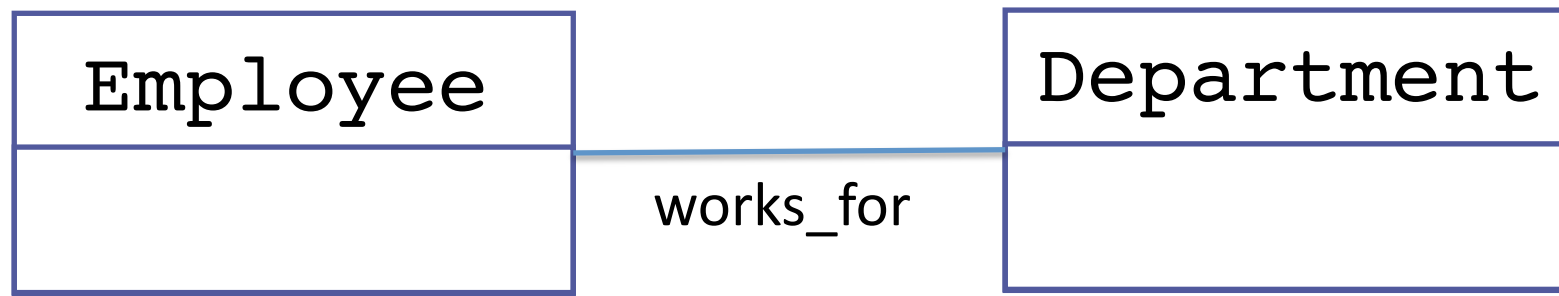
- Relationships between instances of classes



Associations : a tricky notation

- Any association is specified by three things
 1. Its name
 2. The cardinality constraints of the class elements
 3. The role of the classes in the association (optional)

Associations : Name (1)




- The name is a label placed in the middle of the edge

Associations : Cardinality (2)




- This means that one instance of A participates in the association with **N** elements of the other class

Associations : Cardinality (2)




This means that an employee works for exactly **1** department

Associations : Cardinality (2)



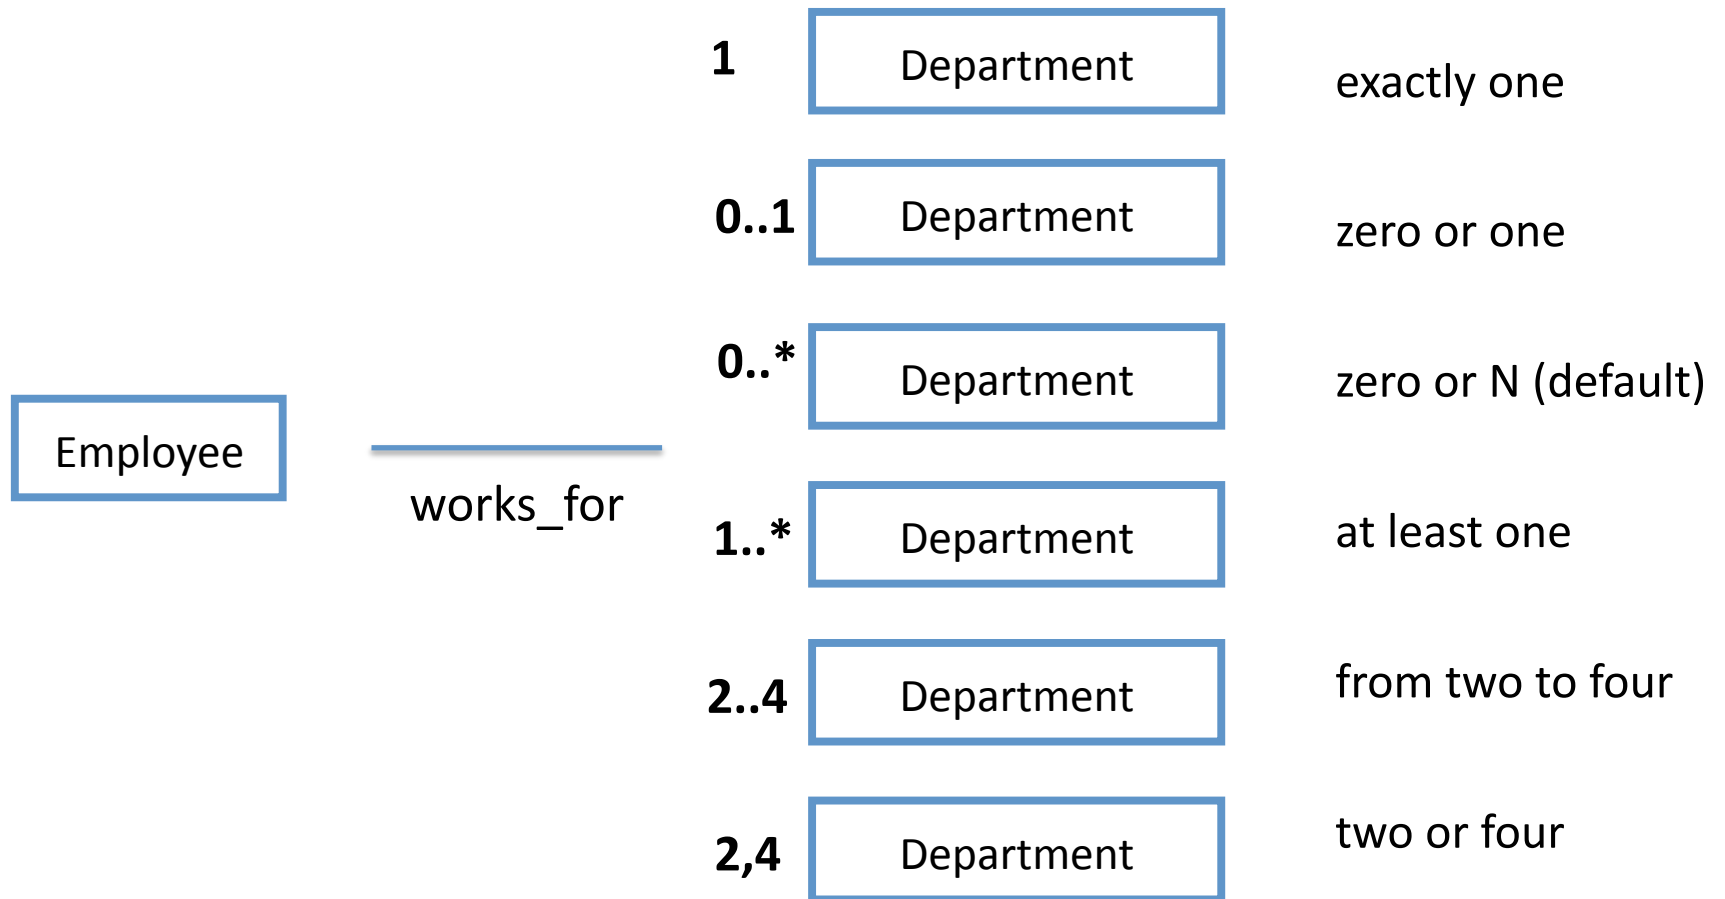
This means that an employee can work for more than **1** department (but at least one)

Associations : Cardinality (2)




This means that a department has at least one employee, with no upper-limit.

Cardinality Specification

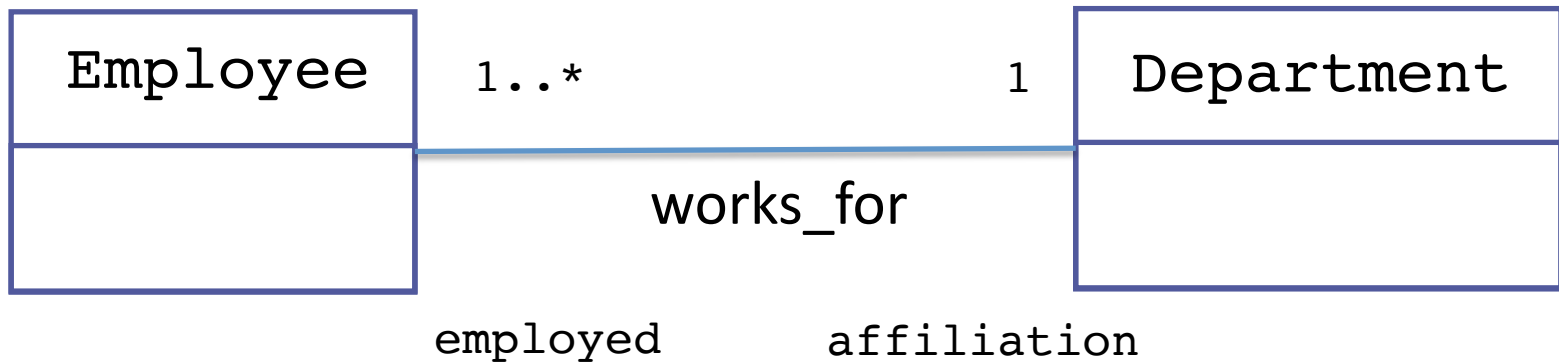


Associations : Roles (2)

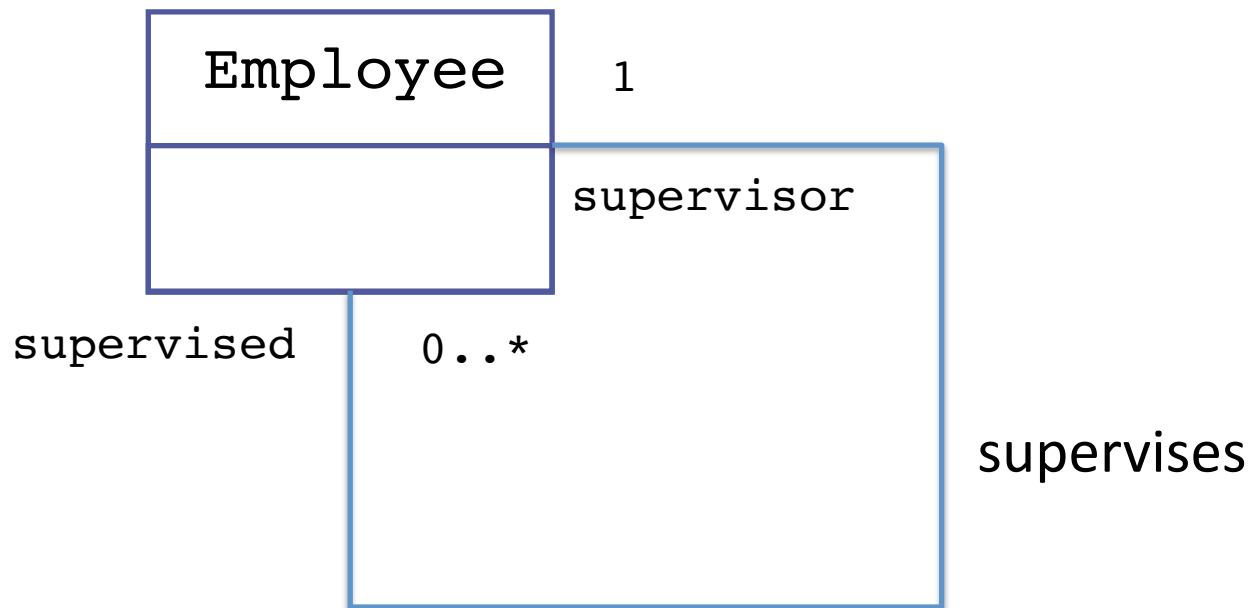


- This means that in the association an instance of A plays **role_A**

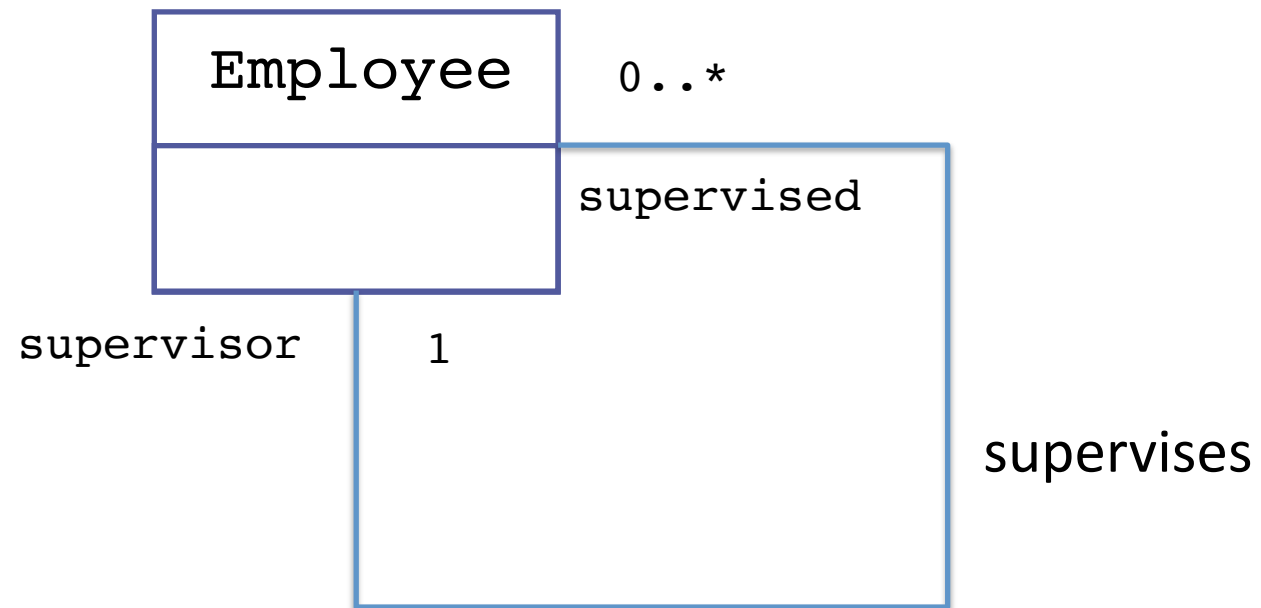
Putting everything together



Reflexive Associations

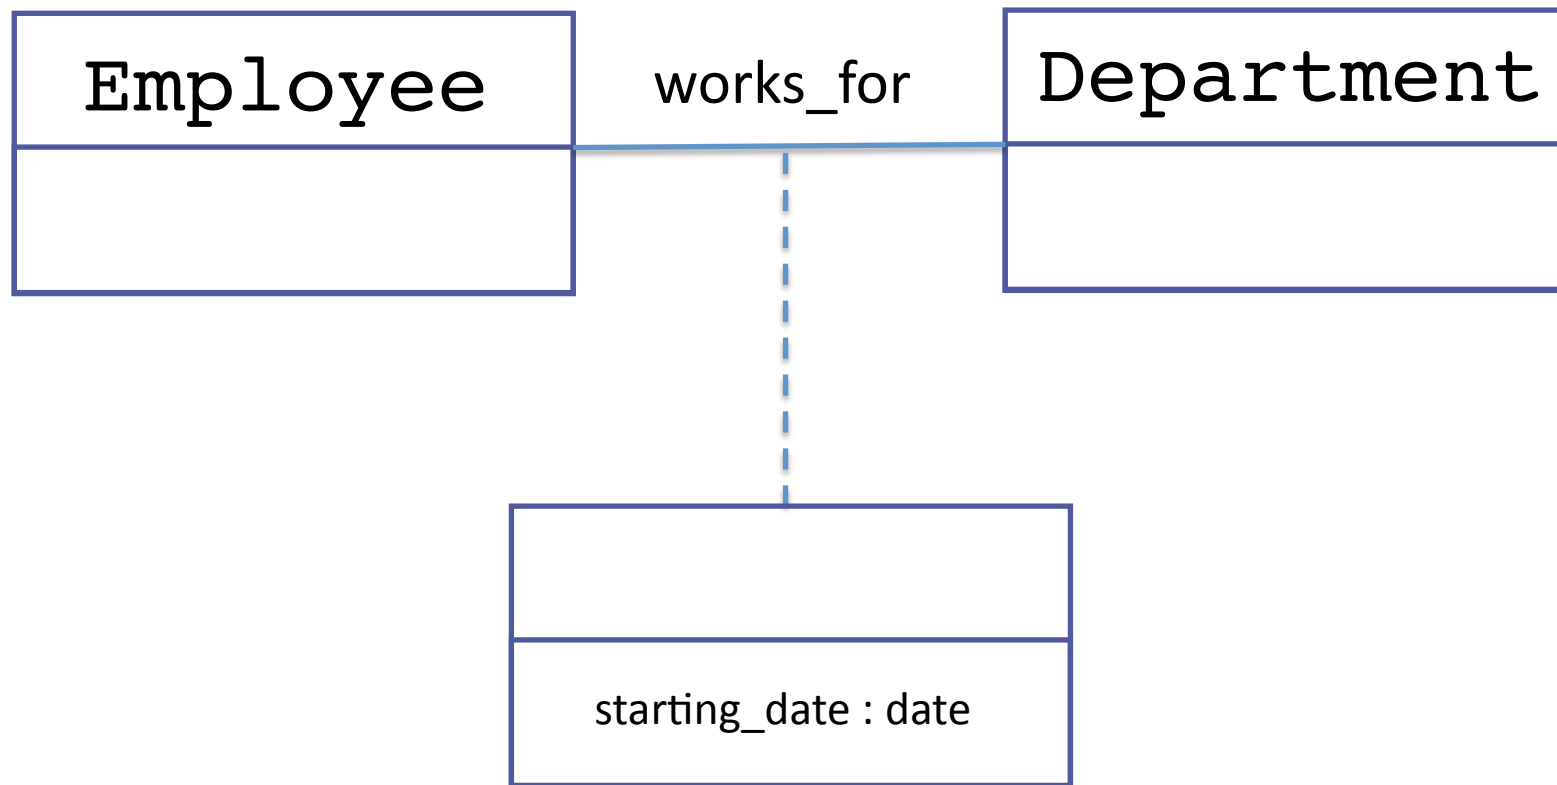


Equivalent Formulation



Notation for Attributes in Associations

[UML] section 4 [BD-G] chapter XVII section 2

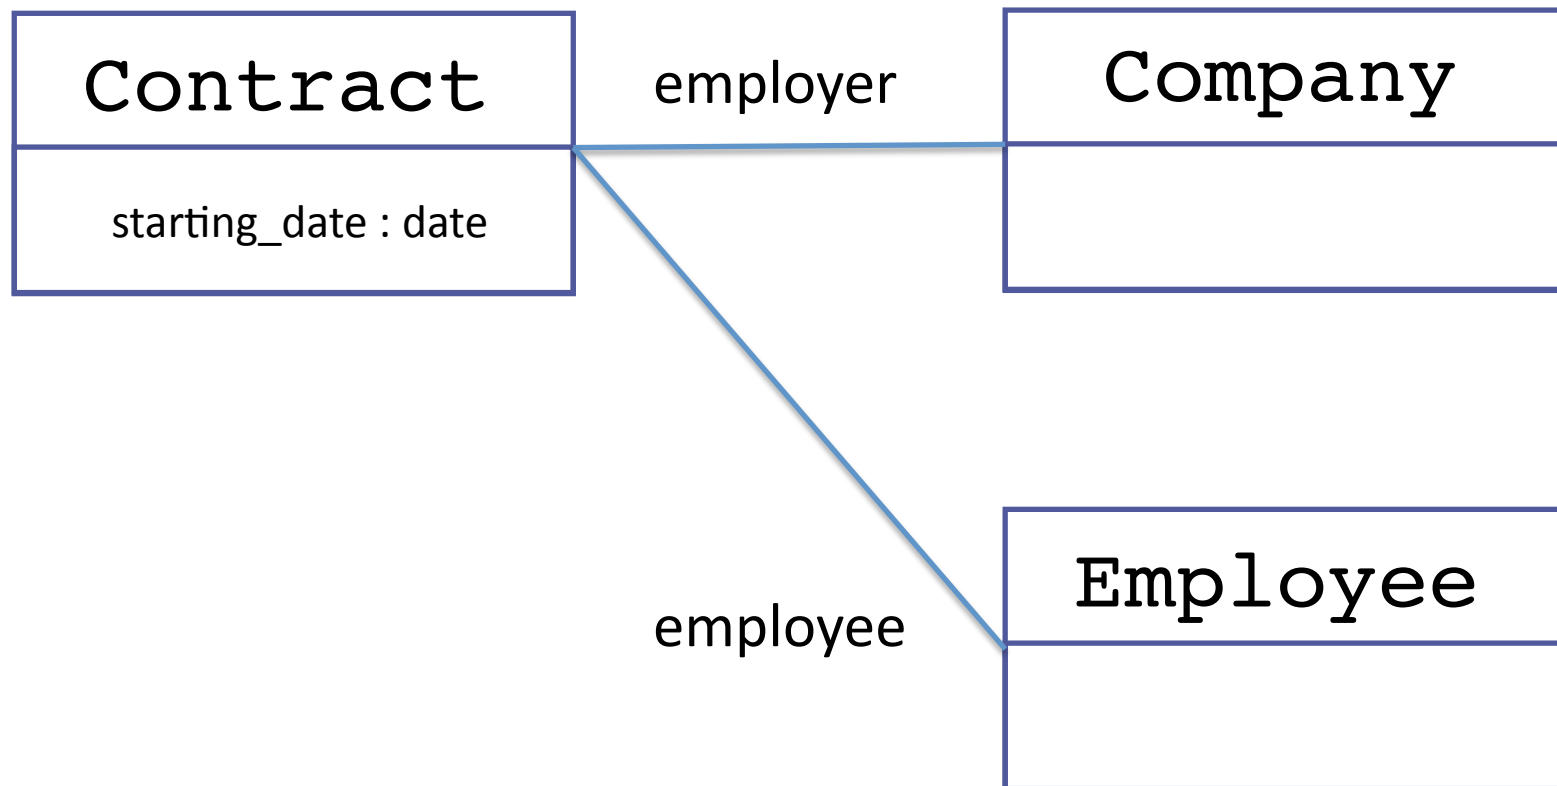


Why Associations can be tricky

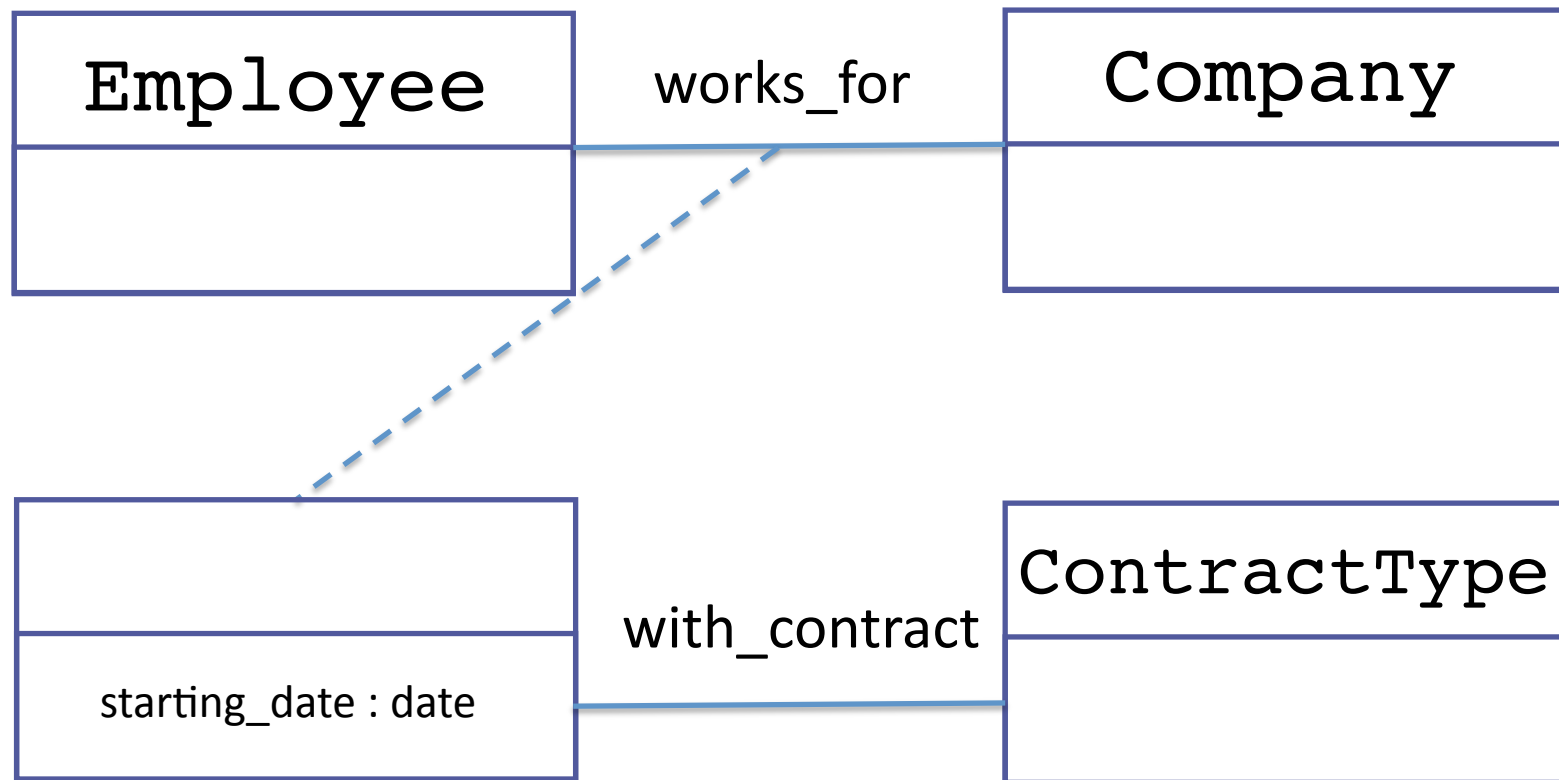
Often, a relationship between two entities combines:

- ***association*** features
 - the fact that two or more things are linked
- **representation** features
 - the details about this association
- Ex: a working contract can be seen as a relationship (an association between an employee and a company) or as an entity (representation of a legal concept)
 - This duality is the source of most design problems !
 - Recognize your own modelling choices !

Why Associations can be Tricky

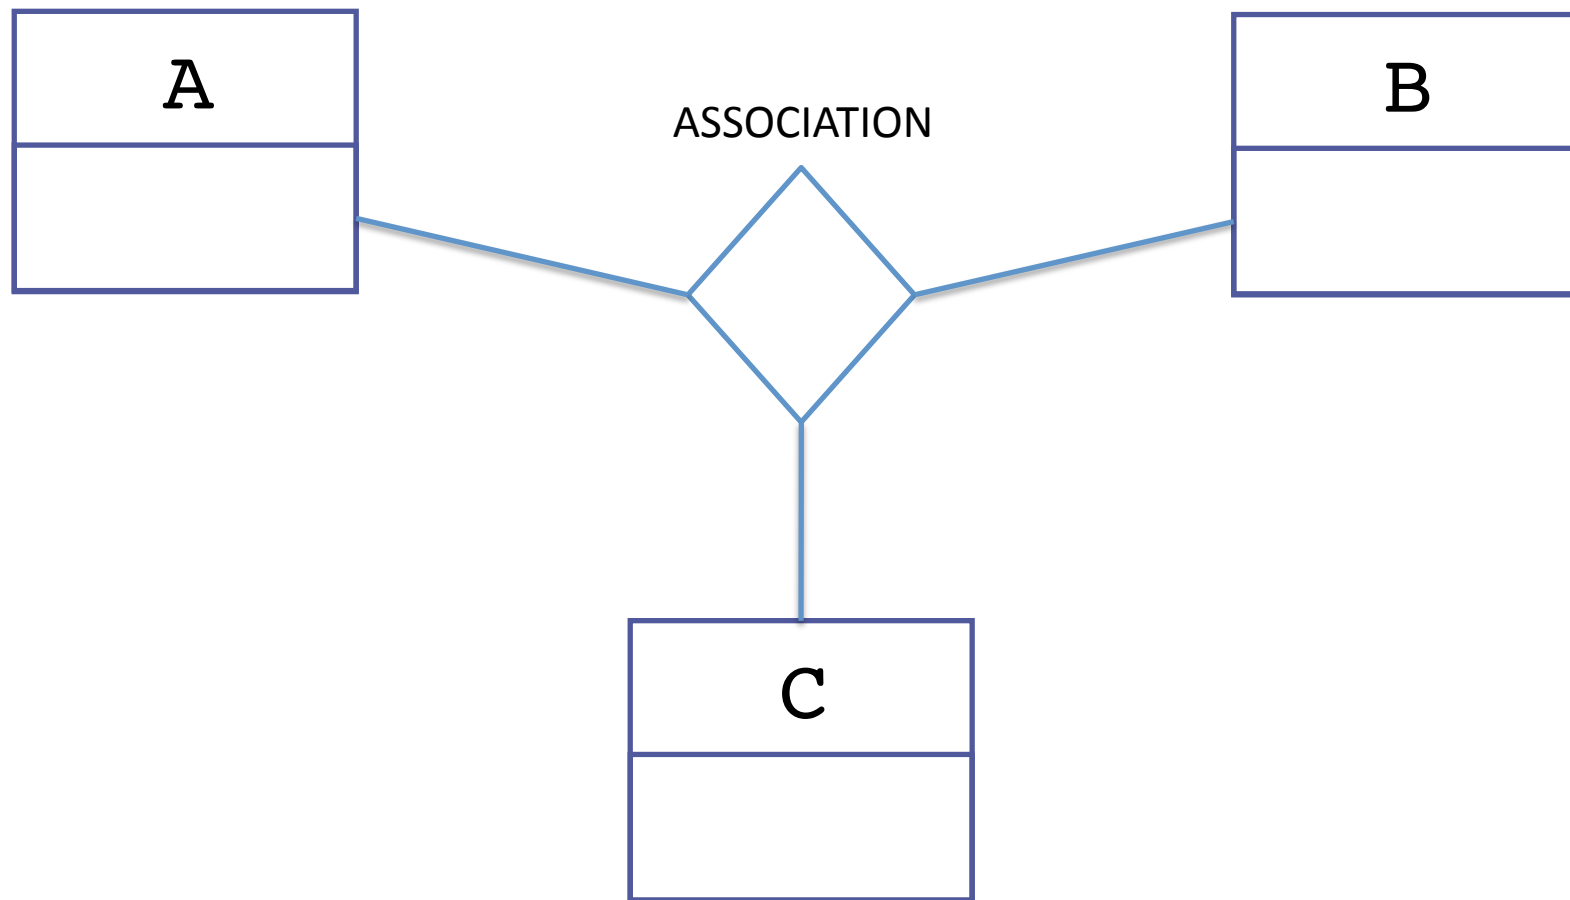


Associations Participating in Other Associations



Notation for 3-ary Associations

[UML] section 4.1 [UML2] section [3.3.4]



Summing Up

UML for Relational Databases


[UML] section 4 [BD-G] chapter XVII section 2

- Basic UML
 - Classes, binary associations, n-ary associations
 - Can model relational databases
- The construct of the language we have seen so far are enough to model relational databases

SubClass


[UML] section 5 [BD-G] chapter XVII section 2.2

- A subset of the instances of a Class
 - ex. every employee is a person
- A subclass inherits all superclass properties or redefines them.



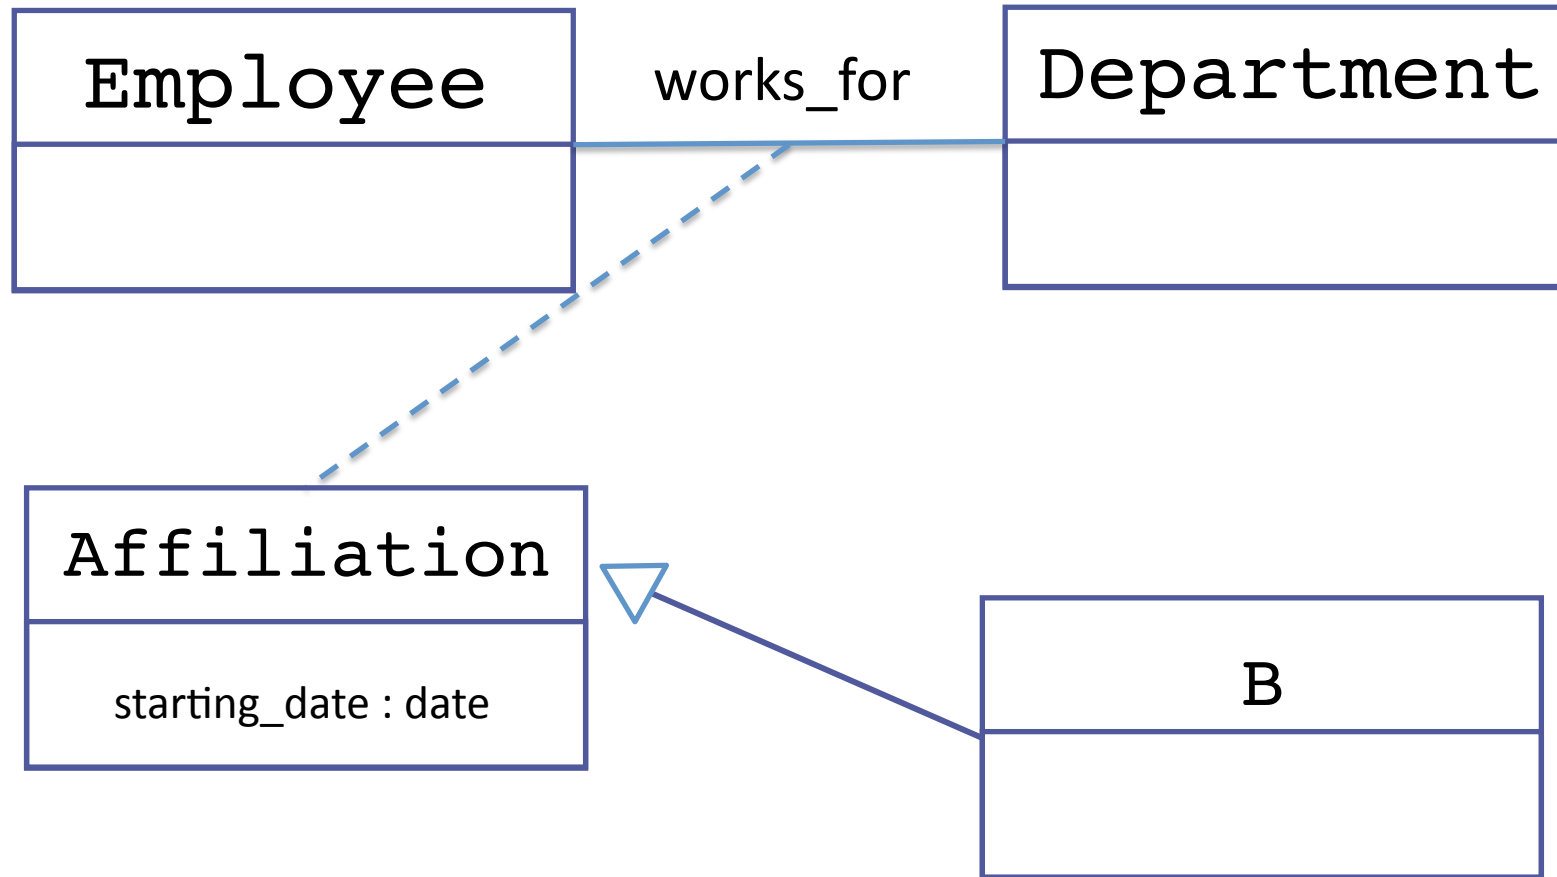
SubClass = Generalization/Specialization

[UML] section 5

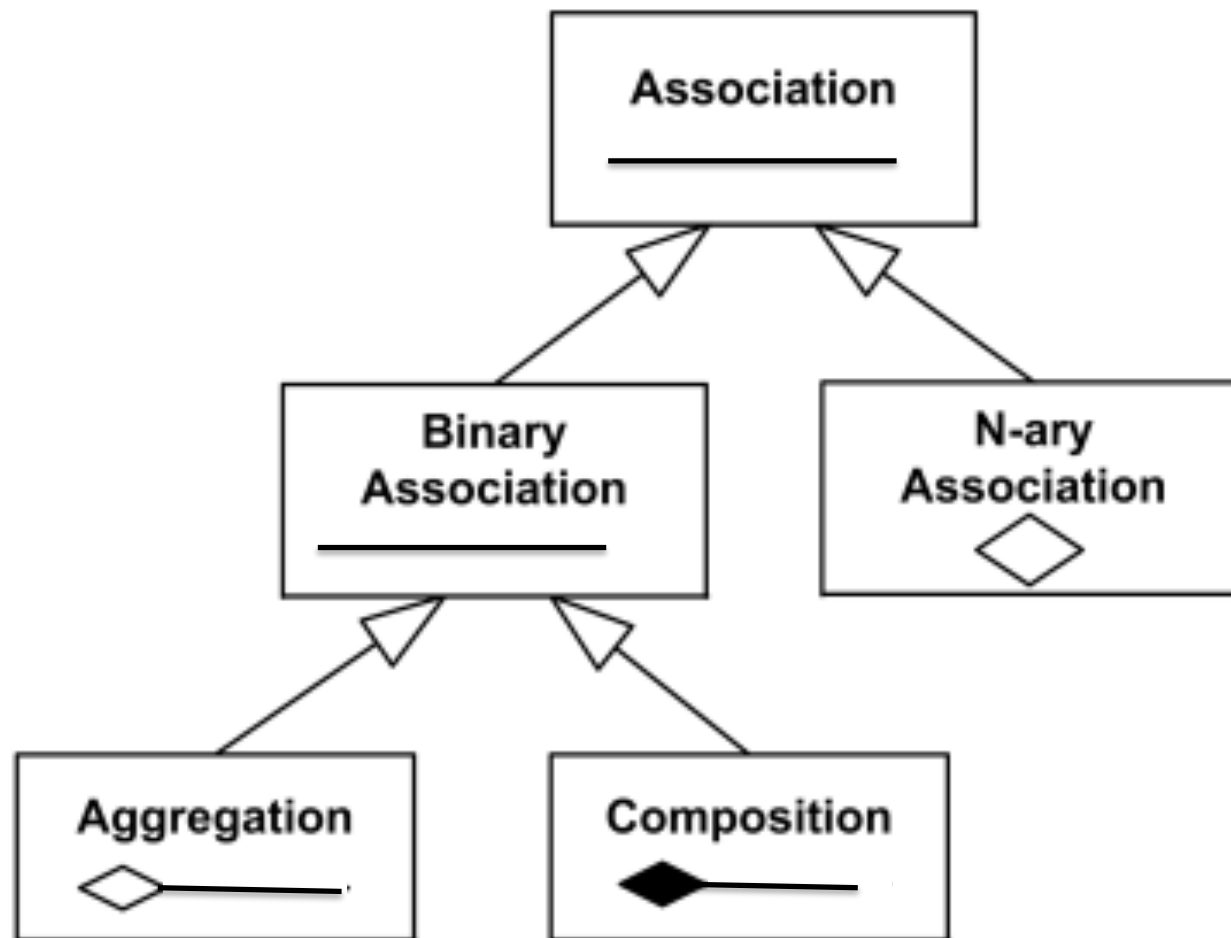


Association modeled by a Class

[UML] section 4.3



- Consequence : such class can have proper attributes, operations and also other associations




Part-Of (binary) Association

- “A forest is made of trees”

Aggregation

[BD-G] chapter XVII section 2.2

- Consequence 1 : a tree can exist even without a forest




Part-Of (binary) Association

- “A folder is made of files”

Composition

[BD-G] chapter XVII section 2.2


- Consequence 1 : a file cannot exists without a folder



Aggregation


[BD-G] chapter XVII section 2.2

- Consequence 1 : a tree can exist even without a forest




Composition

- Consequence 1 : a file cannot exist without a folder




Aggregation

- Consequence 2 : a tree can be part of both a forest and a park



Composition

- Consequence 1 : a file cannot exists without a folder



Summing Up

- Instances/Links/Operations
- Subclasses
- Aggregation and Composition